

# **R Programming and Data Analysis**

## **Intermediate R Programming**

# Introduction

1. Numerical tools
  - a. Optimization
  - b. Integration
  - c. Root finding
  - d. Solving systems of equations
2. Debugging
3. Reading and Writing Data
  - a. Saving and loading workspaces
  - b. CSV files
  - c. HDF5 files
  - d. SAS files
  - e. Databases
  - f. Web services
  - g. Web scraping
  - h. Big datasets
4. Data Manipulation with dplyr

# Numerical Tools

# Numerical Tools

**... Demonstration ...**  
(See `numerical.Rmd`)

# Debugging

# Debugging

- Debugging is an integral part of writing nontrivial programs.
- It is rare to write a perfect program on the first attempt.
- When running a new program, we might see
  - a. errors or warnings.
  - b. an obviously incorrect result (e.g. NA or a negative probability).
  - c. a subtly incorrect result.
  - d. a result whose correctness we are not certain of.
  - e. problems only for certain inputs or random draws.
- Some basic debugging techniques can help us track down the sources of these issues.

# Debugging

- Develop code “interactively”.
  - a. Write a few lines of code, run them, and examine the output.
  - b. Prevents some obvious errors, and quickly get a working program.
- Use print statements to report important values.
  - a. Especially useful for longer programs that cannot be run interactively.
  - b. For example, to see how an optimization is working, we can print the result before returning it.

```
f <- function(x) {  
  fx <- 1 - sum(x^2)  
  cat("x = (", paste(round(x, 4), collapse = ","), "),  
      fx =", fx, "\n")  
  return(fx)  
}
```

```
> x0 <- rep(0, 5)  
> res <- optim(x0, f, control = list(fnscale = -1))  
...  
x = ( 0.005,0.051,0.011,0.011,0.011 ), fx = 0.997011  
x = ( 0.022,0.0244,-0.0916,0.0484,0.0484 ), fx = 0.985845  
x = ( 0.0055,0.0061,0.0521,0.0121,0.0121 ), fx = 0.9969253  
...
```

# Debugging

- Sometimes it is helpful to leave print statements in program after debugging is complete.
- These functions are useful for basic logging

```
printf <- function (msg, ...) {  
  cat(sprintf(msg, ...))  
}  
  
logger <- function (msg, ...) {  
  sys.time <- as.character(Sys.time())  
  cat(sys.time, "-", sprintf(msg, ...))  
}
```

```
> printf("Convergence Status: %d\n", res$convergence)  
Convergence Status: 0  
> logger("Starting %d reps of MCMC\n", R)  
2016-07-25 15:55:01 - Starting 1000 reps of MCMC
```

- The log4r package is a bit more sophisticated. Supports multiple logging levels (INFO, WARN, ERROR, etc).



# Debugging

- R has an interactive debugger to step through running programs.
- Can ask R to start debugger when a particular function is called in the current session. This can be any function in R, not only ours.

```
> debug(optim)
> x0 <- rep(0, 5)
> res <- optim(x0, f, control = list(fnscale = -1))
debugging in: optim(x0, f, control = list(fnscale = -1))
debug: {
... [optim function contents are shown] ...
}
```

- Now program is paused and we can use regular R commands.

```
Browse[2]> ls()
[1] "control" "fn" "gr" "hessian" "lower" "method" "par" "upper"
Browse[2]> print(control)
$fnscale
[1] -1

Browse[2]>
```

- Can step to the next line, the next breakpoint, or stop debugger.
- Changes made to workspace may be discarded after exiting debugger.
- Use `undebbug(optim)` to stop watching `optim` calls.

# Debugging

- Another way to invoke the debugger is to put a browser call in your program. R starts the debugger when it encounters this statement.

```
f <- function(x) {  
  z <- t(x) %% x  
  browser()  
  return(z)  
}
```

```
> x <- c(1,2,3)  
> f(x)  
Called from: f(x)  
Browse[1]> ls()  
[1] "x" "z"  
Browse[1]> x  
[1] 1 2 3  
Browse[1]> z  
      [,1]  
[1,]    14  
Browse[1]> Q
```

- For more information about debugging in R, see [www.stats.uwo.ca/faculty/murdoch/software/debuggingR](http://www.stats.uwo.ca/faculty/murdoch/software/debuggingR).

# Reading and Writing Data

# Objects and the Workspace

- The entities that R creates and manipulates are known as *objects*.
- These may be variables, arrays of numbers, character strings, functions, or more general structures built from such components.
- The collection of objects currently stored in R is called the *workspace*.
- The workspace can be saved to disk, and loaded back into R in a new or existing session.
- Workspace does not store packages that were loaded — we have to reload them ourself.

# Objects and the Workspace

- The commands `object` or `ls` can be used to display the objects currently loaded in R.
- Use the function `rm` to remove objects from your workspace.

```
> x <- rnorm(5)
> x
[1] -0.8287666  0.8377261 -0.3695485  1.3661922  1.8287291
> y <- x + 10
> y
[1]  9.171233 10.837726  9.630451 11.366192 11.828729
> objects()
[1] "x" "y"
> ls()
[1] "x" "y"
> rm(list = ls(all = TRUE))
> ls()
character(0)
```

Note that `character(0)` represents a string vector of length zero.

# Saving the Workspace

- R designates a directory on the computer to be the “current working directory”. This is where output files will be written by default.

```
> getwd()  
[1] "/path/to/here"
```

- To get/set the working directory with `getwd/setwd`

```
> setwd("/path/to/here")  
> getwd()  
[1] "/path/to/here"
```

- When we exit a session, R asks if we wish to save the workspace.
  - ▶ If we say “yes”, a binary file `.RData` will be created; this contains all the objects in our workspace (therefore file might be large).
  - ▶ A text file `.Rhistory` may be also be created; contains our history of commands.
  - ▶ Next time we launch R from that directory, our workspace will return to the same state.
  - ▶ We can also load the state files manually.

# Saving the Workspace

We can save the workspace at any time (not just when quitting), and using any filename we wish, using `save.image`.

```
> x <- c(1,2,3)
> A <- matrix(1:9, nrow = 3, ncol = 3)
> B <- diag(10)
> ls()
[1] "A" "B" "x"
> save.image(file = "myworkspace.Rdata")
```

We can also save specific objects from the workspace using `save`.

```
> ls()
character(0)
> x <- 10
> y <- 11
> z <- 12
> save(list = c("x","y"), file = "myvariables.Rdata")
```

# Loading the Workspace

Saved R data can be loaded at any time using `load`. Be aware that objects in your current environment may be overwritten.

```
> x <- 55
> y <- 77
> load("myworkspace.Rdata")
> ls()
[1] "A" "B" "x" "y"
> x
[1] 1 2 3
> y
[1] 77
```



# CSV Files

- Recall the C02 dataset.

```
> C02
  Plant      Type Treatment conc uptake
1   Qn1    Quebec nonchilled   95   16.0
2   Qn1    Quebec nonchilled  175   30.4
...
83  Mc3 Mississippi   chilled  675   18.9
84  Mc3 Mississippi   chilled 1000   19.9
```

- Write it to a CSV file.

```
> write.table(C02, file = "C02.csv", sep = ",")
> getwd()
[1] "/path/to/file"
```

- Check contents of the file.

```
[username@localhost ~]$ cat /path/to/file/C02.csv
"Plant","Type","Treatment","conc","uptake"
"1","Qn1","Quebec","nonchilled",95,16
"2","Qn1","Quebec","nonchilled",175,30.4
...
"83","Mc3","Mississippi","chilled",675,18.9
"84","Mc3","Mississippi","chilled",1000,19.9
```

# CSV Files

- Read the file into R

```
> dat <- read.table("/path/to/file/C02.csv", sep = ",")
> print(dat)
```

	Plant	Type	Treatment	conc	uptake
1	Qn1	Quebec	nonchilled	95	16.0
2	Qn1	Quebec	nonchilled	175	30.4
...					
83	Mc3	Mississippi	chilled	675	18.9
84	Mc3	Mississippi	chilled	1000	19.9

# Data Input

## Description

Reads a file in table format and creates a data frame from it, with cases corresponding to lines and variables to fields in the file.

## Usage

```
read.table(file, header = FALSE, sep = "", quote = "\"",  
           dec = ".", numerals = c("allow.loss", "warn.loss", "no.loss"),  
           row.names, col.names, as.is = !stringsAsFactors,  
           na.strings = "NA", colClasses = NA, nrows = -1,  
           skip = 0, check.names = TRUE, fill = !blank.lines.skip,  
           strip.white = FALSE, blank.lines.skip = TRUE,  
           comment.char = "#",  
           allowEscapes = FALSE, flush = FALSE,  
           stringsAsFactors = default.stringsAsFactors(),  
           fileEncoding = "", encoding = "unknown", text, skipNul = FALSE)
```

```
read.csv(file, header = TRUE, sep = ",", quote = "\"",  
         dec = ".", fill = TRUE, comment.char = "", ...)
```

```
read.csv2(file, header = TRUE, sep = ";", quote = "\"",  
          dec = ",", fill = TRUE, comment.char = "", ...)
```

```
read.delim(file, header = TRUE, sep = "\t", quote = "\"",  
           dec = ".", fill = TRUE, comment.char = "", ...)
```

# CSV Files

- The readr package provides more sophisticated file parsing tools. Often faster than the usual `read.table`, `read.csv`, etc.
- We will generate a large CSV to demonstrate.

```
library(readr)

n <- 5000000
y <- sample(c("YES", "NO"), size = n, replace = TRUE, prob = c(0.1,
  0.9))
DF <- data.frame(
  x = rnorm(n),
  y = y,
  z = rpois(n, 10)
)
write.table(DF, file = "mydata.dat", sep = ",", row.names = FALSE)
```

- This produces a ~124 MB file.

# CSV Files

```
> system.time(dat1 <- read.csv("mydata.dat"))
   user  system elapsed 
31.166   8.792  46.799 
> head(dat1, 3)
      x      y      z
1 2.4157666 NO     5
2 -0.1859725 NO     8
3 -0.3424828 YES    7
```

```
> system.time(dat2 <- read_csv("mydata.dat"))
Parsed with column specification:
cols(
  x = col_double(),
  y = col_character(),
  z = col_integer()
)
=====| 100% 123 MB
   user  system elapsed 
 7.890   1.702  10.677 
> head(dat2, 3)
# A tibble: 6 x 3
      x      y      z
  <dbl> <chr> <int>
1 2.4157666 NO     5
2 -0.1859725 NO     8
3 -0.3424828 YES    7
```

# HDF5 Files

- HDF5 ([www.hdfgroup.org/HDF5](http://www.hdfgroup.org/HDF5)) flexible library for storing and managing data.
- Portable file format with interfaces in C/C++, Fortran, Python, R, and others.
- An HDF5 file has a hierarchical format like a filesystem.
  1. *Groups* are like directories.
  2. *Datasets* are like files, but have a well-defined structure.
  3. *Attributes* are metadata which can be attached to groups and datasets.
- rhdf5 is an R package for manipulating HDF5 files.

```
source("https://bioconductor.org/biocLite.R")  
biocLite("rhdf5")
```

# HDF5 Files

```
library(rhdf5)
h5createFile("mydata.h5")

y <- mtcars$mpg
X <- model.matrix(~ cyl + disp + hp, data = mtcars)
h5createGroup("mydata.h5", "mtcars")
h5write(mtcars, "mydata.h5", "mtcars/mtcars")
h5write(y, "mydata.h5", "mtcars/y")
h5write(X, "mydata.h5", "mtcars/X")

y <- CO2$conc
X <- model.matrix(~ Plant + Type + Treatment + uptake, data = CO2)
h5createGroup("mydata.h5", "CO2")
h5write(mtcars, "mydata.h5", "CO2/CO2")
h5write(y, "mydata.h5", "CO2/y")
h5write(X, "mydata.h5", "CO2/X")
```

```
> h5ls("mydata.h5")
  group  name      otype  dclass      dim
0      /    CO2    H5I_GROUP
1  /CO2    CO2    H5I_DATASET COMPOUND      32
2  /CO2      X    H5I_DATASET   FLOAT 84 x 15
3  /CO2      y    H5I_DATASET   FLOAT      84
4      / mtcars    H5I_GROUP
5 /mtcars      X    H5I_DATASET   FLOAT 32 x 4
6 /mtcars mtcars    H5I_DATASET COMPOUND      32
7 /mtcars      y    H5I_DATASET   FLOAT      32
8 H5close()
```

# HDF5 Files

```
> library(rhdf5)
> h5read("mydata.h5", "mtcars/y")
[1] 21.0 21.0 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19.2 17.8
[12] 16.4 17.3 15.2 10.4 10.4 14.7 32.4 30.4 33.9 21.5 15.5
[23] 15.2 13.3 19.2 27.3 26.0 30.4 15.8 19.7 15.0 21.4

> h5read("mydata.h5", "mtcars/mtcars")
      mpg   cyl  disp    hp  drat    wt   qsec    vs  am   gear   carb
1  21.0     6 160.0   110  3.90  2.620  16.46    0   1     4     4
2  21.0     6 160.0   110  3.90  2.875  17.02    0   1     4     4
...
31 15.0     8 301.0   335  3.54  3.570  14.60    0   1     5     8
32 21.4     4 121.0   109  4.11  2.780  18.60    1   1     4     2

> h5read("mydata.h5", "mtcars/X")
      [,1] [,2] [,3] [,4]
[1,]     1     6 160.0   110
[2,]     1     6 160.0   110
...
[31,]     1     8 301.0   335
[32,]     1     4 121.0   109

> h5read("mydata.h5", "mtcars/X", index=list(1:5,1:3))
      [,1] [,2] [,3]
[1,]     1     6  160
[2,]     1     6  160
[3,]     1     4  108
[4,]     1     6  258
[5,]     1     8  360
```



# SAS Files

- SAS is a commercial software suite which is popular with data analysts, corporations, and government agencies.
- The main format for SAS files is a proprietary `sas7bdat` format.
  1. The `foreign` and `Hmisc` packages can read them, but require SAS installed on the computer.
  2. The `sas7bdat` and `haven` packages can read them without SAS.
- The `XPORT` format is intended to be more interoperable.
  1. Can be read by the `foreign` package.
- The `foreign` package can also read/write data for statistical packages such as Minitab, S, SAS, SPSS, Stata, Systat.

# SAS Files

```
> library(haven)
> url1 <- "http://www.principlesofeconometrics.com/sas/yield.sas7bdat"
> yield <- read_sas(url1)
> class(yield)
[1] "tbl_df"      "tbl"          "data.frame"
> tail(yield)
# A tibble: 6 x 5
  YIELD      T      GROW      GERM      FLOWER
  <dbl> <dbl>   <dbl>   <dbl>   <dbl>
1 1.080829  34 1.354559 1.155823 1.208213
2 1.831250  35 1.005840 0.991480 0.939275
3 1.028031  36 1.288040 0.366613 0.810828
4 1.465217  37 0.828458 1.385182 0.698436
5 1.706897  38 0.772018 0.837972 0.586044
6 1.988593  39 1.052202 0.895763 1.111877
```

```
> url2 <- "http://www.principlesofeconometrics.com/sas/vote.sas7bdat"
> vote <- read_sas(url2)
> tail(vote)
# A tibble: 6 x 7
  STATE      VOTE      INCOME      SCHOOL      URBAN      NORTHEAST      SOUTHEAST
  <chr>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>
1 Vermont      0 12.415    12.5     0.0         1         0
2 Virginia      0 14.579    12.4    65.6         1         0
3 Washington      0 14.962    12.7    71.1         0         0
4 WestVirginia      1 12.007    12.1    36.1         1         0
5 Wisconsin      1 15.064    12.5    63.0         0         0
6 Wyoming      1 12.6      12.6     0.0         0         0
```

# SAS Files

```
> library(sas7bdat)
> yield <- read.sas7bdat(url1)
> at <- attributes(yield)
> names(at)
[1] "names"           "row.names"       "class"           "pkg.version"
[5] "column.info"     "date.created"    "date.modified"   "SAS.release"
[9] "SAS.host"        "OS.version"      "OS.maker"        "OS.name"
[13] "endian"          "winunix"
> at$SAS.host
[1] "WIN"
> at$SAS.release
[1] "9.0000M0"
> at$date.created
[1] "2008-05-13 17:02:32 EDT"
> at$date.modified
[1] "2008-05-13 17:02:32 EDT"
> str(at$column.info[[1]])
List of 11
 $ name   : chr "YIELD"
 $ offset: int 0
 $ length: int 8
 $ type   : chr "numeric"
 $ fhdr   : int 0
 $ foff   : int 0
 $ flen   : int 0
 $ label  : chr "wheat yield, tonnes per hectare"
 $ lhdr   : int 0
 $ loff   : int 36
 $ llen   : int 31
```

# Databases

- R can interact with both SQL databases (DBs) and NoSQL DBs by using appropriate packages.
- SQL DBs
  - ▶ Store data in a highly structured relational DB
  - ▶ Tables are optimized for storage and efficient merging. This usually requires careful planning.
  - ▶ Use SQL (Structured Query Language) to query and modify the DB.
  - ▶ Examples include MySQL, Oracle, PostgreSQL, SQLite, and SQLServer.
- NoSQL DBs
  - ▶ Less structured than relational DBs (e.g. key-value pairs), but more flexible.
  - ▶ Use an alternative query language to SQL.
  - ▶ Examples include MongoDB and Google BigTable.

# Databases

- We will focus on SQL databases.
- The DBI package provides a generic interface to SQL DBs.
- Other packages build on DBI for specific database implementations. For example, RMySQL, ROracle, RPostgreSQL, RSQLite, RSQLServer.
- SQLite
  - ▶ A “self-contained, serverless, zero-configuration, transactional SQL database engine”.
  - ▶ Databases are stored in ordinary files.
  - ▶ Good for local storage in an application (e.g. storing website cookies in Firefox).

# SQL Databases

```
> library(DBI)
> con <- dbConnect(RSQLite::SQLite(), "mydata.sqlite")

> dbListTables(con)
character(0)
> dbWriteTable(con, "mtcars", mtcars)
[1] TRUE
> dbListTables(con)
[1] "mtcars"

> dbListFields(con, "mtcars")
[1] "row_names" "mpg"      "cyl"      "disp"     "hp"       "drat"     "wt"       "qsec"
[9] "vs"        "am"        "gear"     "carb"
> dbReadTable(con, "mtcars")
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
...											
Maserati Bora	15.0	8	301.0	335	3.54	3.570	14.60	0	1	5	8
Volvo 142E	21.4	4	121.0	109	4.11	2.780	18.60	1	1	4	2

# SQL Databases

```
> res <- dbSendQuery(con, "select * from mtcars where cyl = 4")
> qry <- dbFetch(res)
> dbGetInfo(res)$fields
```

	name	Sclass	type	len
1	row_names	character	TEXT	NA
2	mpg	double	REAL	8
3	cyl	double	REAL	8
4	disp	double	REAL	8
5	hp	double	REAL	8
6	drat	double	REAL	8
7	wt	double	REAL	8
8	qsec	double	REAL	8
9	vs	double	REAL	8
10	am	double	REAL	8
11	gear	double	REAL	8
12	carb	double	REAL	8

```
> dbGetRowCount(res)
[1] 11
> qry$row_names
[1] "Datsun 710"      "Merc 240D"      "Merc 230"      "Fiat 128"
[5] "Honda Civic"    "Toyota Corolla" "Toyota Corona" "Fiat X1-9"
[9] "Porsche 914-2"  "Lotus Europa"   "Volvo 142E"
> qry$mpg
[1] 22.8 24.4 22.8 32.4 30.4 33.9 21.5 27.3 26.0 30.4 21.4
> dbClearResult(res)
[1] TRUE
```

# SQL Databases

```
> res <- dbSendQuery(con, "update mtcars set mpg = -22 where cyl = 4")
> dbClearResult(res)
[1] TRUE
> dbReadTable(con, "mtcars")
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	-22.0	4	108.0	93	3.85	2.320	18.61	1	1	4	1
...											
Volvo 142E	-22.0	4	121.0	109	4.11	2.780	18.60	1	1	4	2

```
> res <- dbSendQuery(con, "insert into mtcars
+ (row_names, mpg, cyl, disp, hp, drat, wt, qsec, vs, am, gear, carb)
+ values ('ABCD', 40.0, 4, 100.0, 200, 4.00, 1.5, 20.0, 1, 1, 5, 1)")
> dbClearResult(res)
[1] TRUE
> dbReadTable(con, "mtcars")
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
...											
Volvo 142E	-22.0	4	121.0	109	4.11	2.780	18.60	1	1	4	2
ABCD	40.0	4	100.0	200	4.00	1.500	20.00	1	1	5	1

```
> dbDisconnect(con)
[1] TRUE
```



# Web Services

- Web services are collections of functions that can be called through the web by user programs.
- User programs call functions via Application Programmer Interface (API).
- A simple API, useful for querying, is REST (Representational State Transfer). REST embeds queries in standard (HTTP) web requests.
- JavaScript Object Notation (JSON) is a popular format for returning data.
- The R package `jsonlite` can be used to query a REST service that returns JSON.

# Web Services

- Transport for Finland lets us query real-time positions of its trams.
- Users often have to register for an API key before using web service. No key is required for this one.
- For more information about the service, see <http://digitransit.fi/en/developers/services-and-apis>.
- If we access the URL directly,

```
$ curl http://api.digitransit.fi/realtime/vehicle-positions/v1/hfp/
journey/tram/#
{"hfp/journey/tram/RHKL00401/1007A/2/XXX/2247/undefined
/60;24/29/17/50":{"VP":{"desi":"1007A","dir":"2","oper":"XXX","
veh":"RHKL00401","tst":"2016-08-02T01:42:55.000Z","tsi
":1470102175,"spd":0,"hdg":210,"lat":60.215367,"long
":24.970981,"dl":567,"oday":"XXX","jrn":"XXX","line":"1007A","
start":"2247","stop_index":17}}}
```

# Web Services

- Reading it with jsonlite package,

```
> library(jsonlite)
> url <- "http://api.digitransit.fi/realtime/vehicle-positions/v1/
  hfp/journey/tram/#"
> req <- fromJSON(url)
> str(req)
List of 1
 $ /hfp/journey/tram/RHKL00401/1007A/2/XXX/2247/undefined
  /60;24/29/17/50:List of 1
   ..$ VP:List of 16
    .. ..$ desi      : chr "1007A"
    .. ..$ dir       : chr "2"
    .. ..$ veh       : chr "RHKL00401"
    .. ..$ tst       : chr "2016-08-02T01:39:24.000Z"
    .. ..$ tsi       : int 1470101964
    .. ..$ hdg       : int 210
   ...
    .. ..$ lat       : num 60.2
    .. ..$ long      : num 25
    .. ..$ dl        : int 567
    .. ..$ line      : chr "1007A"
    .. ..$ start     : chr "2247"
    .. ..$ stop_index: int 17
```

# Web Scraping

- Data on the web is often prepared for viewing (rendered in HTML) rather than for analysis (e.g. a CSV or Excel file).
- “Web scraping” is the process of writing a script to extract the data from the HTML.
- The `rvest` package supports web scraping in R.

# Web Scraping

Adam Jones Stats - Baltimore Orioles - ESPN - Mozilla Firefox

Adam Jones Stats - Balti... x

espn.go.com/mlb/player/stats/\_id/28513/adam-jones

Search

Adam Jones

#10 CF Bats: R, Throws: R Baltimore Orioles

Birth Date August 1, 1985 (Age: 30)  
Birthplace San Diego, CA  
Experience 10 years  
College None  
HT/WT 6-2, 215 lbs.

2016 Season

AVG	HR	RBI	OBP
.262	17	53	.308

Career

.277	213	712	.318
------	-----	-----	------

Go to Adam Jones

Player Profile Stats **Split**s Game Log Bat vs Pitch Hot Zones Videos Photos Shop

## Adam Jones Stats

Batting Fielding

### CAREER BATTING STATISTICS

YEAR	TEAM	GP	AB	R	H	2B	3B	HR	RBI	BB	SO	SB	CS	AVG	OBP	SLG	OPS	WAR
2006	SEA	32	74	6	16	4	0	1	8	2	22	3	1	.216	.237	.311	.548	0.0
2007	SEA	41	65	16	16	2	1	2	4	4	21	2	1	.246	.300	.400	.700	0.8
2008	BAL	132	477	61	129	21	7	9	57	23	108	10	3	.270	.311	.400	.711	1.7
2009	BAL	119	473	83	131	22	3	19	70	36	93	10	4	.277	.335	.457	.792	2.4
2010	BAL	149	581	76	165	25	5	19	69	23	119	7	7	.284	.325	.442	.767	2.1
2011	BAL	151	567	68	159	26	2	25	83	29	113	12	4	.280	.319	.466	.785	2.8
2012	BAL	162	648	103	186	39	3	32	82	34	126	16	7	.287	.334	.505	.839	3.4
2013	BAL	160	653	100	186	35	1	33	108	25	136	14	3	.285	.318	.493	.811	4.4
2014	BAL	159	644	88	181	30	2	29	96	19	133	7	1	.281	.311	.469	.780	4.9
2015	BAL	137	546	74	147	25	3	27	82	24	102	3	1	.269	.308	.474	.782	3.3
2016	BAL	87	355	57	93	14	0	17	53	24	69	1	0	.262	.308	.445	.753	--
Total	Total	1329	5083	732	1409	243	27	213	712	243	1042	85	32	.277	.318	.461	.779	--
Season Averages		120.0	462.1	66.5	128.1	22.1	2.5	19.4	64.7	22.1	94.7	7.7	2.9	.277	.318	.461	.779	--

### POSTSEASON BATTING STATISTICS

SEASON	GP	AB	R	H	2B	3B	HR	RBI	BB	SO	SB	CS	AVG	OBP	SLG	OPS
--------	----	----	---	---	----	----	----	-----	----	----	----	----	-----	-----	-----	-----

# Web Scraping

```
library(rvest)

my.html <- read_html("http://espn.go.com/mlb/player/stats/_/id/28513/
adam-jones")

all.stats <- my.html %>%
  html_table(head = TRUE)

stats <- all.stats[[2]]
View(stats)
```

The %>% operator is defined in the magrittr package.

- Feeds data into a function.
- Many of them can be strung together.
- Behaves like the “pipe” operator on the Linux command line.

# Web Scraping

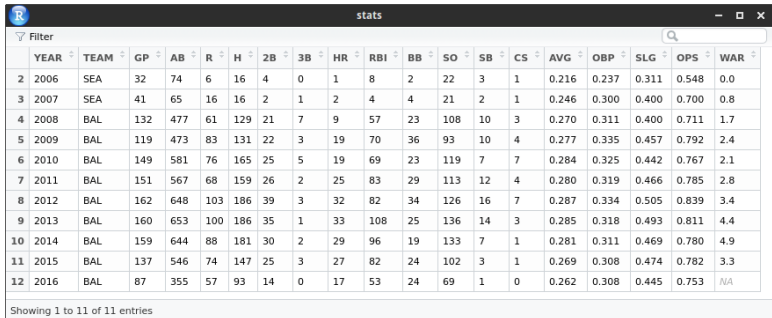
stats								
Filter								
	CAREER BATTING STATISTICS	CAREER BATTING STATISTICS	CAREER BATTING STATISTICS	CAREER BATTING STATISTICS	CAREER BATTING STATISTICS	CAREER BATTING STATISTICS	CAREER BATTING STATISTICS	CAREER BATTING STATISTICS
1	YEAR	TEAM	GP	AB	R	H	2B	3B
2	2006	SEA	32	74	6	16	4	0
3	2007	SEA	41	65	16	16	2	1
4	2008	BAL	132	477	61	129	21	7
5	2009	BAL	119	473	83	131	22	3
6	2010	BAL	149	581	76	165	25	5
7	2011	BAL	151	567	68	159	26	2
8	2012	BAL	162	648	103	186	39	3
9	2013	BAL	160	653	100	186	35	1
10	2014	BAL	159	644	88	181	30	2
11	2015	BAL	137	546	74	147	25	3
12	2016	BAL	87	355	57	93	14	0
13	Total	Total	1329	5083	732	1409	243	27
14	Season Averages	Season Averages	120.0	462.1	66.5	128.1	22.1	2.5
Showing 1 to 14 of 14 entries								

# Web Scraping

```
# Fix the header
colnames(stats) <- stats[1,]

# Remove the extra first row, and the row of totals and averages
stats <- stats[-c(1,13,14),]

# Change columns 3, ..., k to numeric
k <- ncol(stats)
for (j in 3:k) {
  stats[,j] <- as.numeric(stats[,j])
}
```



	YEAR	TEAM	GP	AB	R	H	2B	3B	HR	RBI	BB	SO	SB	CS	AVG	OBP	SLG	OPS	WAR
2	2006	SEA	32	74	6	16	4	0	1	8	2	22	3	1	0.216	0.237	0.311	0.548	0.0
3	2007	SEA	41	65	16	16	2	1	2	4	4	21	2	1	0.246	0.300	0.400	0.700	0.8
4	2008	BAL	132	477	61	129	21	7	9	57	23	108	10	3	0.270	0.311	0.400	0.711	1.7
5	2009	BAL	119	473	83	131	22	3	19	70	36	93	10	4	0.277	0.335	0.457	0.792	2.4
6	2010	BAL	149	581	76	165	25	5	19	69	23	119	7	7	0.284	0.325	0.442	0.767	2.1
7	2011	BAL	151	567	68	159	26	2	25	83	29	113	12	4	0.280	0.319	0.466	0.785	2.8
8	2012	BAL	162	648	103	186	39	3	32	82	34	126	16	7	0.287	0.334	0.505	0.839	3.4
9	2013	BAL	160	653	100	186	35	1	33	108	25	136	14	3	0.285	0.318	0.493	0.811	4.4
10	2014	BAL	159	644	88	181	30	2	29	96	19	133	7	1	0.281	0.311	0.469	0.780	4.9
11	2015	BAL	137	546	74	147	25	3	27	82	24	102	3	1	0.269	0.308	0.474	0.782	3.3
12	2016	BAL	87	355	57	93	14	0	17	53	24	69	1	0	0.262	0.308	0.445	0.753	NA

Showing 1 to 11 of 11 entries



# Big Datasets

- Base R expects all objects (vectors, matrices, data frames) to be completely loaded into memory.
- Modern datasets may be too large to fit into memory.
- The `bigmemory` and `ff` packages store matrices and data frames on disk, but allow them to be accessed somewhat like regular R objects.
  1. The `ff` project: <http://ff.r-forge.r-project.org>
  2. The `bigmemory` project: <http://www.bigmemory.org>

# Data Manipulation with dplyr

# Tidyverse

- “an opinionated collection of R packages designed for data science. All packages share an underlying philosophy and common APIs.”

**R's biggest challenge** is  
that most **R** users are  
not programmers.



Hadley Wickham

- Tidyverse includes:
  - a. ggplot2 - grammar for plotting.
  - b. dplyr - grammar for data manipulation. (\*\*\*)
  - c. readr, readxl, haven - reading data from files.
  - d. magrittr - provides the pipe operator (%>%).
- Many tidyverse packages are discussed in the book *R for Data Science* (?), which is available online.

# Data Manipulation with dplyr

**... Demonstration ...**  
(See `dplyr.Rmd`)

# References I